

Homework 3 Solutions

Barras Simon <simon.barras@epfl.ch>

Zachary Doll <zachary.doll@epfl.ch>

1 Problem 1 (Frugal Vertex Coloring)

Let $G = (V, E)$ be a graph with maximum degree Δ . A vertex coloring $\chi : V \rightarrow C$ is called β -frugal if for every vertex v , no color appears more than β times in its neighborhood $N(v)$. That is,

$$|\{u \in N(v) : \chi(u) = c\}| \leq \beta \quad \text{for all } v \in V, c \in C.$$

A coloring is *proper* if $\chi(u) \neq \chi(v)$ for all edges $\{u, v\} \in E$.

Show that for any constant integer $\beta \geq 1$, there exists a β -frugal coloring of G using $Q = O(\Delta^{1+1/\beta})$ colors.

In fact, a stronger statement is true: there is a coloring that is both proper and β -frugal. But we only require to prove the above weaker statement.

1.1 Solution

Proof. Consider a color set C with cardinality $|C| := Q$. Let us color each vertex of $G = (V, E)$ uniformly at random:

$$\chi(v) \stackrel{\text{i.i.d.}}{\sim} \mathcal{U}\{Q\}, \quad \forall v \in V$$

For each vertex $v \in V$ all colours $c \in C$, and any subset $S \subseteq N(v)$ of size $|S| = \beta + 1$, define the bad event:

$$B(v, c, S) = \{\text{all vertices in } S \text{ have color } c\}$$

If such a bad event exists, then by definition, the coloring χ is not β -frugal over G .

From the uniform color assignment, we have that the probability p of the bad event occurring for a fixed tuple (v, c, S) is given by:

$$\mathbb{P}[B(v, c, S)] = Q^{-(\beta+1)} \triangleq p$$

Let us now find an upper-bound for the number of events d^* on which $B(u, c, S)$ depends, for a fixed vertex $u \in S$:

1. Because each color is assigned independently for each vertex, the event $B(u, c, S)$ only depends on the colors in S . In other words, for the events $B(u, c, S)$ and $B(w, \tilde{c}, \tilde{S})$ to be dependent, it must be the case that $u \in \tilde{S} \subseteq N(w)$, i.e. w and u must be neighbors.

By hypothesis, we know that $\delta(v) \leq \Delta$ for all $v \in V$, so there are at most Δ choices for an initial vertex.

2. We have at most $\binom{\Delta-1}{\beta}$ choices for the remaining β vertices in $\tilde{S} \subseteq N(w) \setminus \{u\}$.
3. We have Q choices for the color $\tilde{c} \in C$.

Thus,

$$d^* \leq \Delta \binom{\Delta-1}{\beta} Q - 1$$

Where we subtracted 1 for the event itself. Union-bounding over the fixed cardinality $|S| = \beta + 1$ yields an upper-bound to the total number of events d that any non-fixed event depends on:

$$\begin{aligned} d &\leq (\beta + 1)d^* \\ &\leq Q(\beta + 1)\Delta \binom{\Delta-1}{\beta} - 1 \\ &\leq Q(\beta + 1)\Delta \left(\frac{e\Delta}{\beta}\right)^\beta - 1 \\ &= Q(\beta + 1) \left(\frac{e}{\beta}\right)^\beta \Delta^{\beta+1} - 1 \stackrel{\Delta}{=} QC_\beta \Delta^{\beta+1} - 1 \end{aligned}$$

Where $C_\beta = (\beta + 1) \left(\frac{e}{\beta}\right)^\beta$, and where we used stirling's bound in the third line:

$$k! \geq \left(\frac{k}{e}\right)^k \implies \frac{n^k}{k!} \leq \frac{n^k}{(k/e)^k} = \left(\frac{en}{k}\right)^k$$

The Local Lovász Lemma states that, if $ep(d + 1) \leq 1$, then with positive probability, no bad event occurs, meaning a β -frugal coloring exists. Thus, plugging in the computed value for p and the above bound for d into the LLL inequality, it suffices to have

$$eQ^{-(\beta+1)} \cdot QC_\beta \Delta^{\beta+1} \leq 1 \implies eC_\beta Q^{-\beta} \Delta^{\beta+1} \leq 1$$

Finally, solving for Q gives:

$$Q^\beta \geq eC_\beta \Delta^{\beta+1} \implies Q \geq (eC_\beta)^{1/\beta} \Delta^{1+1/\beta}$$

And thus

$$Q = O\left(\Delta^{1+1/\beta}\right)$$

as required. □

2 Problem 2 (Concentration for Euclidean MST)

Let X_1, \dots, X_n be n points chosen independently and uniformly at random from the unit square $[0, 1]^2$. Let $L(X_1, \dots, X_n)$ denote the total length of the Minimum Spanning Tree (MST) on these points, using Euclidean distances.

Let $\mu = \mathbb{E}[L(X_1, \dots, X_n)]$ be the expected length of the MST. Prove that for any $\epsilon > 0$, the probability of deviating from the mean by ϵn is exponentially small in n . Specifically, show that:

$$\Pr(|L - \mu| \geq \epsilon n) \leq 2 \exp\left(-\frac{\epsilon^2 n}{25}\right)$$

Hint: You may use the following fact without proof.

Fact: Any Euclidean MST on points in the 2D plane (using the L_2 norm) has a maximum vertex degree of at most 5.

2.1 Solution

We aim to prove that L is a Lipschitz function, and then apply McDiarmid's inequality to obtain the desired concentration bound.

Proof. We claim that L is a Lipschitz function. That is, for any two configurations $X = (X_1, \dots, X_n)$ and $X' = (X_1, \dots, X'_i, \dots, X_n)$ that differ only in one coordinate i , the following holds:

$$|L(X) - L(X')| \leq c$$

for some constant c .

Let T be the minimum spanning tree (MST) of X , and let T' be the tree obtained from T by replacing the vertex X_i with X'_i . Although T' may not be the MST of X' , it is still a valid spanning tree on the modified set of points. let T'^* denote the MST of X' , then:

$$\text{len}(T'^*) \leq \text{len}(T')$$

Since all points lie within the unit square $[0, 1]^2$, the maximum Euclidean distance between any two point is

$$\sqrt{1^2 + 1^2} = \sqrt{2}$$

Given the fact that in planar MST, each vertex has maximum degree of 5. Therefore, when a single point X_i is moved to X'_i , at most five edges are affected, each by changing by at most $\sqrt{2}$ in length. Hence

$$|\text{len}(T) - \text{len}(T')| \leq 5\sqrt{2}$$

Since $\text{len}(T'^*) \leq \text{len}(T')$, it follows that:

$$|L(X) - L(X')| \leq |\text{len}(T) - \text{len}(T'^*)| \leq 5\sqrt{2}$$

Thus, L is c -Lipschitz with $c = 5\sqrt{2}$.

We can now apply McDiarmid's inequality, which states that if L is c_i -Lipschitz in each coordinate, then:

$$\Pr [|L - \mu| \geq t] \leq 2 \exp\left(\frac{-2t^2}{\sum_{i=1}^n c_i^2}\right)$$

Since each $c_i = 5\sqrt{2}$, we have

$$\sum_{i=1}^n c_i^2 = n \cdot (5\sqrt{2})^2 = 50n$$

Setting $t = \varepsilon n$, we obtain:

$$\Pr [|L - \mu| \geq \varepsilon n] \leq 2 \exp\left(\frac{-2\varepsilon^2 n^2}{50n}\right) = 2 \exp\left(\frac{-\varepsilon^2 n}{25}\right)$$

□

3 Problem 3 (A randomized algorithm for k -SAT)

Consider a satisfiable k -CNF Φ on n variables. One try of the algorithm: start at uniform $x_0 \in \{0, 1\}^n$; for T steps $t = 0, 1, \dots, T-1$, if x_t satisfies Φ return x_t , else pick an unsatisfied clause C , choose a uniform random literal $\ell \in C$ and flip its variable to obtain x_{t+1} from x_t . If no solution within T steps, restart. Fix a satisfying assignment x^* and let $D_t = \|x_t - x^*\|_1$.

- (a) Show that whenever $D_t > 0$, $\Pr[D_{t+1} = D_t - 1 | x_t] \geq 1/k$.
- (b) If $D_0 = d$, prove $\Pr[\text{hit 0 within } d \text{ steps}] \geq (1/k)^d$ (via d consecutive decreases).
- (c) For x_0 uniform, $D_0 \sim \text{Bin}(n, 1/2)$. Show $\Pr[\text{success in one try}] \geq \left(\frac{k+1}{2k}\right)^n$.
- (d) Argue that $T = n$ suffices to capture the event in (b), and conclude the expected time $\tilde{O}\left((\frac{2k}{k+1})^n\right)$; specialize to $k = 3$ as $\tilde{O}\left((\frac{3}{2})^n\right)$.

Remark (Schöning's bound). If in (b) you instead bound $\Pr[\text{ever hit 0} | D_0 = d] \geq (1/(k-1))^d$ using a biased random-walk/gambler's-ruin argument with step -1 w.p. $1/k$ and $+1$ w.p. $1 - 1/k$, then averaging as in (c) yields per-try success $\left(\frac{k}{2(k-1)}\right)^n$ and expected time $\tilde{O}\left((2 - \frac{2}{k})^n\right)$ (e.g., $\tilde{O}\left((\frac{4}{3})^n\right)$ for 3-SAT).

3.1 Solution

3.1.1 Part A

Proof. Pick an arbitrary assignment x_t with distance $D_t > 0$ from a satisfying x^* , and assume x_t does not satisfy Φ (in which case the algorithm would terminate).

Because x^* satisfies Φ , every clause C contains at least one literal $\ell^* \in C$ set to true under the assignment x^* . Conversely, as C is unsatisfied under x_t , every literal of C must be false in x_t . In particular, we have that $\ell^* \in C$ is false in x_t . Because the algorithm picks a literal $\ell \in C$ u.a.r. among the k possible choices, and at least one of them (ℓ^*) would decrease the Hamming distance by 1, it follows that

$$\Pr[D_{t+1} = D_t - 1 | x_t] \geq \frac{1}{k}$$

□

3.1.2 Part B

Proof. We assume the satisfying assignment x^* is unique, as per clarifications on the exercise.

Consider the event that each of the first d steps decrease the distance by exactly 1. From part (a), we have that

$$\Pr \left[\bigcap_{t=0}^{d-1} \{D_{t+1} = D_t - 1\} \right] = \prod_{t=0}^{d-1} \Pr[D_{t+1} = D_t - 1 | x_t] \geq \left(\frac{1}{k}\right)^d$$

Thus:

$$\Pr[\text{hit 0 within } d \text{ steps} | D_0 = d] \geq \left(\frac{1}{k}\right)^d$$

□

3.1.3 Part C

Proof. the result from part (b) gives us

$$\Pr[\text{success in one try} | D_0 = d] \geq \left(\frac{1}{k}\right)^d$$

From which it immediately follows that

$$\begin{aligned} \Pr[\text{success in one try}] &\geq \mathbb{E} \left[\left(\frac{1}{k}\right)^{D_0} \right] = \prod_{i=1}^n \mathbb{E} \left[\left(\frac{1}{k}\right)^Z \right], \quad Z \sim \text{Ber}(1/2) \\ &= \prod_{i=1}^n \left(\frac{1}{2} + \frac{1}{2} \cdot \frac{1}{k} \right) \\ &= \left(\frac{k+1}{2k} \right)^n \end{aligned}$$

Where we simply used the MGF of $D_0 \sim \text{Bin}(n, 1/2)$.

□

3.1.4 Part D

The event from (b), namely:

$$D_0 = d \implies \Pr[\text{hit 0 within } d \text{ steps}] \geq (1/k)^d$$

is fully contained in the first d steps. Because we consider the K-SAT problem over n variables, it is clear that the maximal possible initial distance is $D_0 = n$, i.e. when all variables from $x_0 \sim \mathcal{U}\{0,1\}^n$ differs from the satisfying assignment x^* . Thus, $D_0 = d \leq n$, and choosing $T = n$ suffices to capture the event fully.

From part (c), we know that the probability of success in one try is:

$$P \geq \left(\frac{k+1}{2k} \right)^n$$

Therefore the expected number of necessary tries until success is

$$\mathbb{E}[\text{necessary tries until success}] = O(1/P) = O\left(\left(\frac{2k}{k+1}\right)^n\right)$$

Each try costs $T = n$ time, and thus:

$$\mathbb{E}[\text{time}] = O\left(n\left(\frac{2k}{k+1}\right)^n\right) = \tilde{O}\left(\left(\frac{2k}{k+1}\right)^n\right)$$

4 Problem 4 (The Long(est) Path Home)

Given a graph $G = (V, E)$, you want to find long simple paths in the graph in polynomial time.

- (a) (Algorithm 1: Dead easy.) Show that you can find a path of length k (if such a path exists) in time $n\Delta^k$, where Δ is the maximum degree of G .
- (b) (Easy.) If the graph were directed and acyclic (i.e., a DAG), then show that you can deterministically find the longest path in G in time $O(m + n)$. Here, and in general, $m = |E|$ and $n = |V|$.
- (c) (Algorithm 2:) Consider running the following procedure n times, and outputting the longest path found in these n tries.

Take a random permutation of the vertices, and direct each edge from the lower endpoint to the higher endpoint to create a DAG \vec{G} . Find a longest path in \vec{G} .

Show that for $k = c \frac{\log n}{\log \log n}$ for some constant $c > 0$, Algorithm 2 will find a path of length k (if it exists) with probability at least $1/2$.

- (d) Now, consider a slight extension of this idea. Suppose you have a graph G , and you color the vertices using k colors (neighbors need not have different color). A path is called *polychromatic* if has $\ell \leq k$ vertices, and all the ℓ vertices have different colors.
 - i. Show that you can find a polychromatic path of length k in time that is $\text{poly}(n, k)2^k$. (So, this is polynomial time for $k = O(\log n)$).
 - ii. (Algorithm 3:) Consider running the following procedure n times, and outputting the longest path found in these n tries.

Take a random coloring of the vertices using k colors, and find the polychromatic path of length at most k in G .

Show that for $k = c \log n$ for some constant $c > 0$, Algorithm 3 will find a path of length k (if it exists) with probability at least $1/2$. (Hint: Use Stirling's approximation.)

4.1 Solution

4.1.1 Part A

We aim to show that the brute-force algorithm for finding the longest path of length at most k in a graph $G = (V, E)$ runs in $O(n\Delta^k)$, where $n = |V|$ is the number of vertices and Δ is the maximum degree of any vertex in G .

Proof. The high-level algorithm (Algorithm 1) enumerates all vertices as potential starting points for the longest path. For each vertex $v \in V$, it calls a recursive procedure that explores all possible simple paths starting from v with length at most k . The outer loop thus contributes as factor of $O(n)$.

Algorithm 1 Find longest path in the graph

```

procedure LONGEST-PATH( $G = (V, E), k$ )
     $P \leftarrow \emptyset$ 
    for  $v \in V$  do
         $Q \leftarrow \text{Longest-Path-Vertex}(G, v, \emptyset, k)$ 
         $P \leftarrow \max(P, Q)$ 
    end for
    return  $P$ 
end procedure
```

Algorithm 2 describes the recursive procedure that computes the longest path starting from a fixed vertex v . At each recursive step, the algorithm considers all neighbors u of v (excluding the predecessor vertex p) and recursively explores all paths of remaining length $k - 1$ starting from u .

Since each vertex has at most Δ neighbors, and the recursion proceeds to depth k , the total number of recursive calls can be upper-bounded by $O(\Delta^k)$. This corresponds to exploring all possible paths of length up to k starting from a given vertex.

Algorithm 2 Find longest path of a vertex

```

procedure LONGEST-PATH-VERTEX( $G = (V, E), v, p, k$ )
     $P \leftarrow \emptyset$ 
    if  $k = 0$  then
        return  $P$ 
    end if
    for  $u \in V$  where  $\{u, v\} \in E$  and  $u \neq p$  do
         $Q \leftarrow \text{Longest-Path-Vertex}(G, u, v, k - 1)$ 
         $P \leftarrow \max(P, Q)$ 
    end for
    return  $P$ 
end procedure
```

Combining both procedures, the overall time complexity is obtained by multiplying the outer $O(n)$ loop with the recursive exploration cost $O(\Delta^k)$. Therefore, the total running time of the algorithm is

$$O(n\Delta^k)$$

□

4.1.2 Part B

We aim to show that if the graph $G = (V, E)$ is directed and acyclic, there exist an algorithm to find the longest path in time $O(V + E)$.

Proof. Finding the the longest path in a DAG (Directed Acyclic Graph) can achieved efficiently using a DFS (Depth-First Search) combined with dynamic programming. This approach ensures that each vertex is processed only once, avoiding redundant computations and achieving the desired linear-time complexity.

The high-level procedure (Algorithm 3) iterates over all vertices, treating each as potential starting points for the longest path. For every vertex $v \in V$, it invokes a recursive routines that computes and memorizes the longest path starting from v . The outer loop thus contributes as factor of $O(V)$.

Algorithm 3 Find longest path in a DAG

```

procedure LONGEST-PATH-DAG( $G = (V, E), k$ )
     $\mathcal{P} \leftarrow \{\emptyset\}$  for  $|V|$ 
    for  $v \in V$  do
        Longest-Path-Vertex-DAG( $G, \mathcal{P}, v$ )
    end for
    return  $\max(P \in \mathcal{P})$ 
end procedure
```

Algorithm 4 defines the recursive routine improved with dynamic programming that computes the longest path starting from a fixed vertex v . If the longest path for v has already been computed (i.e., $\mathcal{P}[v] \neq \emptyset$), the result is returned immediately. Otherwise, for each outgoing edge $(v, u) \in E$, the algorithm recursively computes the longest path starting from u and updates $\mathcal{P}[v]$ accordingly.

Algorithm 4 Find longest path of a vertex in DAG

```

procedure LONGEST-PATH-VERTEX-DAG( $G = (V, E), \mathcal{P}, v$ )
    if  $\mathcal{P}[v] \neq \emptyset$  then return
    end if
    for  $e = \{v, u\} \in E$  do
        Longest-Path-Vertex-DAG( $G, \mathcal{P}, u$ )
         $\mathcal{P}[v] \leftarrow \max(\mathcal{P}[v], \mathcal{P}[u] + e)$ 
    end for
end procedure
```

Each vertex v is visited only once, and during its processing, the algorithm inspects all its outgoing edges. The total amount of work fo all vertices is the total number of edges E .

Thus combining bot procedures yields an algorithm that computes the longest path in a DAG in linear time.

$$O(V + E)$$

□

4.1.3 Part C

We claim that if we assign a random ordering to all vertices of an undirected graph $G = (V, E)$ and orient each edge from the vertex with smaller index to the one with larger index, we obtain a directed acyclic graph (DAG) $\vec{G} = (V, E)$.

Now, consider construction n independent random orientations $\vec{G}_1, \dots, \vec{G}_n$. We want to show that, for

$$k = c \frac{\log n}{\log \log n}$$

with a sufficiently large constant $c > 0$, the probability that at least one of these DAGs contains a path of length k is at least $\frac{1}{2}$:

$$\Pr \left[\exists i, P \subseteq \vec{G}_i, \text{len}(P) = k \right] \geq \frac{1}{2}$$

Proof. Fix a specific path P of length k in G . It consists of $k+1$ vertices. When we randomly permute the vertices, there are $(k+1)!$ possible relative orderings of these vertices, and only one of them produces an increasing order consistent with the directed edges of \vec{G} . Hence

$$\Pr \left[P \subseteq \vec{G} \right] = \frac{1}{(k+1)!} =: p$$

The longer the path, the smaller this probability becomes. Let P^* be a longest path in G . For our family of random orientations, the probability that at least one of the \vec{G}_i contains P^* is

$$\Pr \left[P \in \bigcup_{i=1}^n \vec{G}_i \right] \geq \Pr \left[P^* \in \bigcup_{i=1}^n \vec{G}_i \right] = 1 - \Pr \left[P^* \notin \bigcup_{i=1}^n \vec{G}_i \right] = 1 - (1-p)^n$$

Using the standard exponential bound $(1-p)^n \leq e^{-pn}$.

$$\Pr \left[P^* \in \bigcup_{i=1}^n \vec{G}_i \right] \geq 1 - e^{-pn} = 1 - \exp \left(-n \frac{1}{(k+1)!} \right) = 1 - \exp \left(-n \frac{1}{\left(c \frac{\log n}{\log \log n} + 1 \right)!} \right)$$

Using Stirling's bound

$$(k+1)! \geq \sqrt{2\pi(k+1)} \left(\frac{k+1}{e} \right)^{k+1}$$

we obtain

$$\frac{1}{(k+1)!} \leq \frac{1}{\sqrt{2\pi(k+1)}} \left(\frac{e}{k+1} \right)^{k+1}$$

Substituting $k = c \frac{\log n}{\log \log n}$

$$p \leq \frac{1}{\sqrt{2\pi k}} \left(\frac{e \log \log n}{c \log n} \right)^{c \frac{\log n}{\log \log n}}$$

Hence

$$pn \leq n^{1-c+o(1)}$$

If we choose any $c > 1$, then $pn \rightarrow 0$ as $n \rightarrow \infty$, implying $1 - e^{-pn} \rightarrow 0$. Conversely, if $c < 1$, then $pn \rightarrow \infty$, and

$$1 - e^{-pn} \rightarrow 1$$

Thus, there exist a threshold constant $c^* \approx 1$ such that for $c < c^*$ the probability exceeds $\frac{1}{2}$ for large n . \square

4.1.4 Part D

Consider a graph $G = (V, E)$ where each vertex is assigned one of k colors, such that no two adjacent vertices share the same color. A path P is said to be polychromatic if it contains at most one vertex of each color. Clearly, the length of any such path satisfies $\text{len}(P) \leq k$.

We claim that there exist algorithms running in time $\text{poly}(n, k)2^k$ that finds the longest polychromatic path.

Proof. Consider a dynamic programming (DP) table defined as:

$$\text{DP}[C][v] = \begin{cases} \text{True} & \text{if there exists a path ending at vertex } v \text{ with } C \\ \text{False} & \text{otherwise} \end{cases}$$

Where C is a polychromatic subset.

The DP is initialized for single-vertex paths:

$$\text{DP}[\text{col}(v)][v] = \text{True}, \quad \forall v \in V$$

Then, it is updated recursively:

$$\text{DP}[C][v] = \bigvee_{u \in N(v)} \text{DP}[C \setminus \text{col}(v)][u]$$

Where $N(v)$ denotes the set of neighbors of v .

Each entry in the DP table represents whether a polychromatic path with color subset C ends at vertex v .

Since there are 2^k subsets of colors and n vertices, the table contains $O(n2^k)$ entries. For each vertex, we may check all its incident edges, leading to a total time complexity of:

$$O\left(2^k(n + m)\right) = \text{poly}(n, k)2^k$$

□

By repeating the random coloring of the graph $G = (V, E)$ independently n times to generate graphs $G_1^\bullet, \dots, G_n^\bullet$, we claim that for

$$k = c \log n$$

with constant $c > 0$, the algorithm finds a polychromatic path of length k (if exists) with probability at least $\frac{1}{2}$:

$$\Pr [\exists i, P \subseteq G_i^\bullet, \text{len}(P) = k] \geq \frac{1}{2}$$

Proof. Fix a specific path P of length k in G . The path contains k distinct vertices, each independently assigned one of k colors.

There are k^k possible colorings of the vertices along P , but only those in which all vertices have distinct colors yield a polychromatic path. The number of distinct-color colorings equals $k!$. Thus, for a single random coloring:

$$\Pr [\text{p}(P) \in G^\bullet] = \frac{k!}{k^k} =: p$$

Where $p(P)$ is the polychromatic possible path of P .

For our family of random coloring, the probability that at least one of the G_i^\bullet contains a valid $p(P)$ of length k is

$$\Pr [\exists i : p(P) \in G_i^\bullet] = 1 - (1-p)^n \geq 1 - e^{-pn} = 1 - \exp\left(-n \frac{k!}{k^k}\right)$$

To analyze this expression asymptotically, we apply Stirling's approximation:

$$k! \approx \sqrt{2\pi k} \left(\frac{k}{e}\right)^k$$

Thus,

$$\frac{k!}{k^k} \approx e^{-k} \sqrt{2\pi k}$$

Plugging this back, we get:

$$\Pr [\exists i : p(P) \in G_i^\bullet] \geq 1 - \exp\left(e^{-k} \sqrt{2\pi k}\right)$$

Setting $k = c \log n$

$$ne^{-k} \sqrt{2\pi k} = ne^{-c \log n} \sqrt{2\pi c \log n} = n^{1-c} \sqrt{2\pi c \log n}$$

For this probability to be at least $\frac{1}{2}$, we require the exponent to be at least $\ln 2$. That is, $n^{1-c} \sqrt{\log n} = \Omega(1)$, which holds for any constant $c < 1$.

Therefore, for any constant $c < 1$,

$$k = c \log n \Rightarrow \Pr [\exists i : p(P) \in G_i^\bullet] \geq \frac{1}{2}$$

□